

ON TOUR | NEW YORK



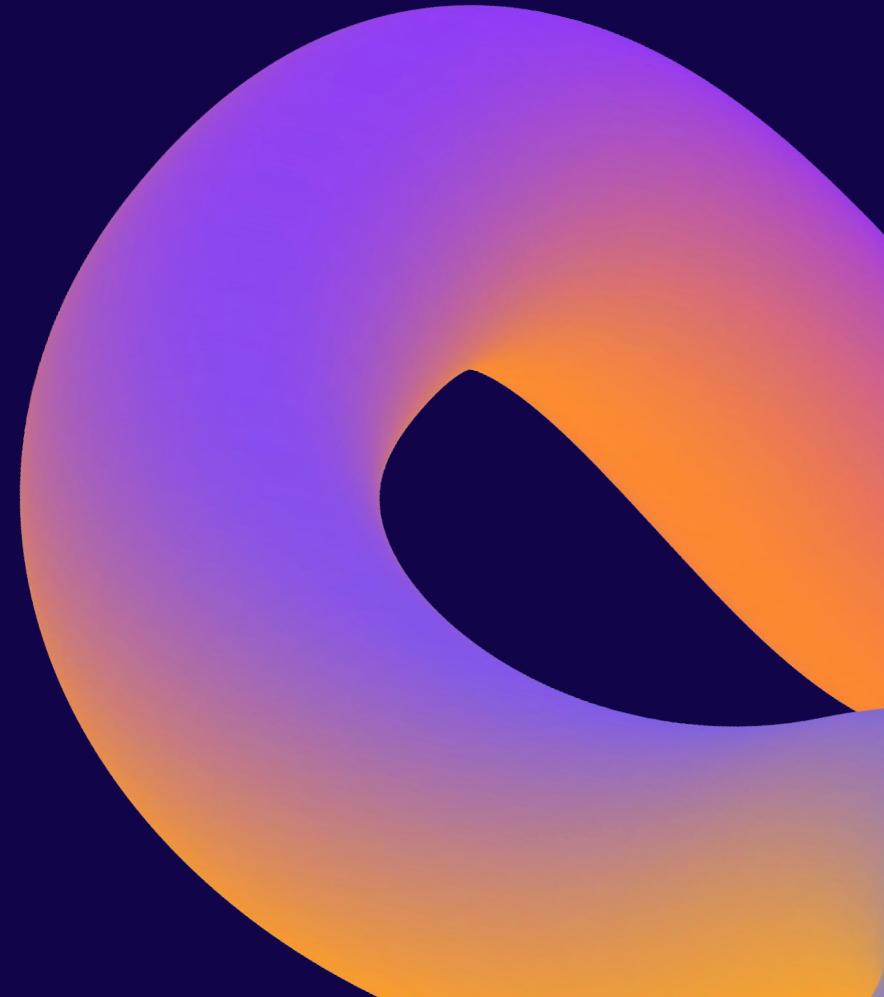
Optimized Locking: Improving SQL Server Transaction Concurrency

Deborah Melkin

she\her

Data Engineer

Advisor360



Deborah Melkin

She/Her

Data Engineer
Advisor360



- . 25 years as a DBA
- . Data Platform Women in Tech (WIT) Virtual UG, Co-leader
- . WITspiration, Co-founder
- . Redgate Community Ambassador
- . Microsoft MVP, Data Platform
- . In my spare time, I can usually be found doing something musical or something geeky with my husband, Andy, and our dog, Sebastian.



@dgmelkin.bsky.social



github.com/DebtheDBA



DebtheDBA.wordpress.com

What's the problem we're trying to solve?

- Locking & Blocking!

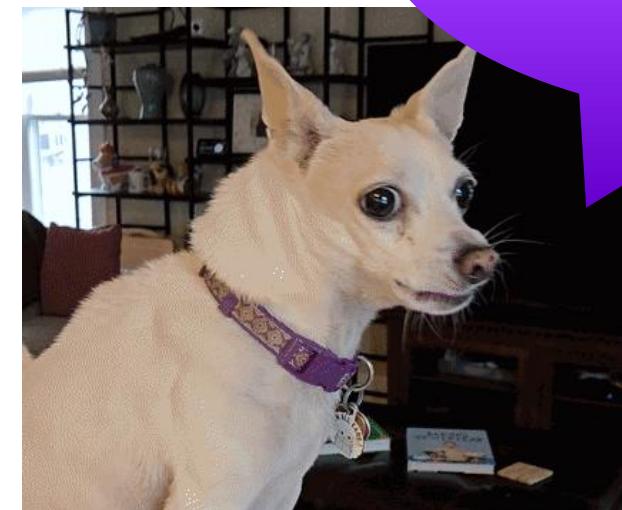


Introducing Optimized Locking

- GA in Azure SQL DB in Feb 2023
- SQL Server 2025 Public Preview
- Writes not blocking Writes

Agenda – How it works

- Components used by Optimized Locking
 - Read Committed Snapshot Isolation Level (RCSI)
 - Accelerated Database Recovery (ADR)
- Lock Escalation
- Optimized Locking
 - Transaction ID Locking
 - Lock After Qualification



Warning:
*** Internals Ahead ***
Level 300

Pessimistic vs Optimistic Locking

- **Pessimistic Locking** – Preventing users from modifying data in a way that affects other users
 - Read Uncommitted
 - Read Committed
 - Serializable
 - Repeatable Read

Pessimistic vs Optimistic Locking

- **Optimistic Locking** – No locks when reading data but data is checked when modifying
 - Read Committed Snapshot Isolation
 - Snapshot Isolation
 - Optimized Locking

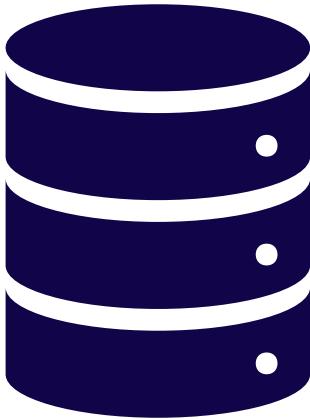
Read Committed Snapshot Isolation Level (RCSI)

- Introduced with SQL Server 2005
- Writes and Reads don't block each other
- Two database options for setting this:
 - SET READ COMMITTED SNAPSHOT ISOLATION – sets default for the database
 - ALLOW SNAPSHOT ISOLATION LEVEL – allows individual sessions\queries to use SNAPSHOT isolation levels even if read committed snapshot isolation is not set.

Read Committed Snapshot Isolation Level (RCSI)

- Every row in the database has a Transaction ID (TID)
- Keeps copies of the previously committed versions of the records in the Version Store in TempDB
- Only reading the committed version of the data

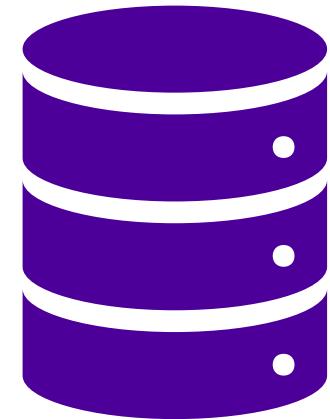
User DB



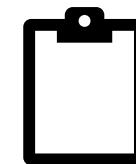
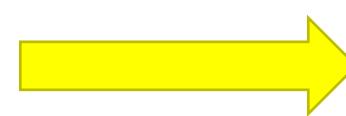
BEGIN TRAN 1A

```
UPDATE Customer  
SET State = 'MA'  
WHERE FirstName = 'Sebastian'
```

tempdb



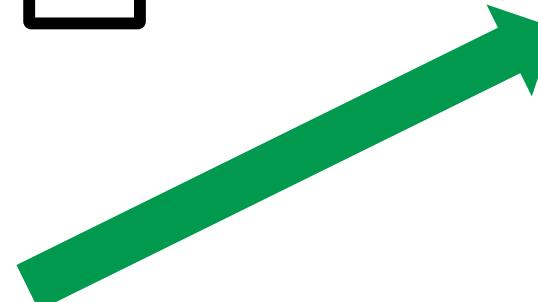
Sebastian	MA	Key	TID 1A
-----------	----	-----	-----------



Version Store			
Sebastian	AZ	TID 1A	1

TID
Transaction Seq No.

```
SELECT FirstName,  
LastName,  
State  
FROM Customer  
WHERE FirstName = 'Sebastian'
```



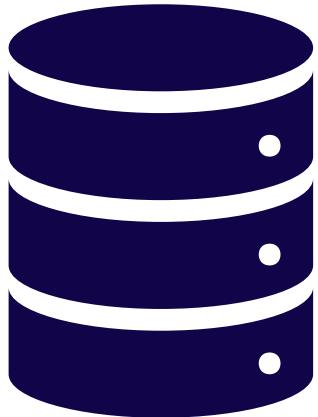
Read Committed
Snapshot Isolation

Accelerated Database Recovery (ADR)

- Introduced in SQL Server 2019
- Default for Azure SQL DB
 - Cannot be turned off
- Changes the way the transactions logs are read to be able to recover from long running transactions or just restore faster
- Uses a Persistent Version Store in Database rather than the Version Store in TempDB

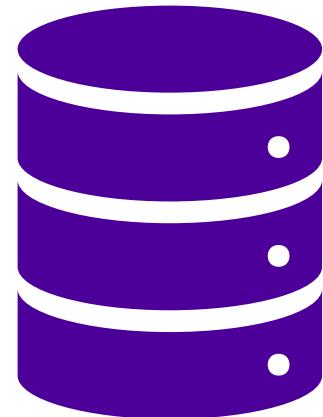
BEGIN TRAN 1A

User DB

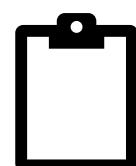


UPDATE Customer
SET State = 'MA'
WHERE FirstName = 'Sebastian'

tempdb



Sebastian	MA	Key	1A
-----------	----	-----	----



Persistent Version Store

Sebastian	AZ	1A	1
-----------	----	----	---

SELECT FirstName,
LastName,
State

FROM Customer
WHERE FirstName = 'Sebastian'

Version Store

Accelerated Database
Recovery

Lock Escalation

- Threshold at which SQL Server will convert lower granularity locks to a higher level in order to manage a small number of locks
 - 5,000 on a table or index for a single statement
 - Each additional 1,250 locks taken in a transaction

Lock Granularity (from Microsoft)

Resource	Description
RID	A row identifier used to lock a single row within a heap.
KEY	A row lock to lock a single row in a B-tree index.
PAGE	An 8 kilobyte (KB) page in a database, such as data or index pages.
EXTENT	A contiguous group of eight pages, such as data or index pages.
HoBT	A heap or B-tree. A lock protecting a B-tree (index) or the heap data pages in a table that doesn't have a clustered index.
TABLE	The entire table, including all data and indexes.
FILE	A database file.
APPLICATION	An application-specified resource.
METADATA	Metadata locks.
ALLOCATION_UNIT	An allocation unit.
DATABASE	The entire database.
XACT	Transaction ID (TID) lock used in Optimized Locking

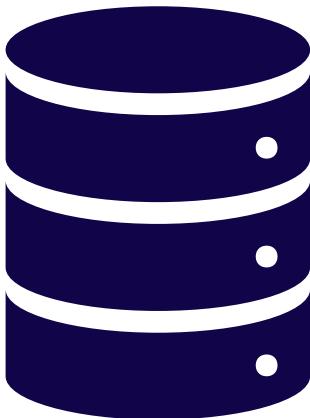
Lock Escalation (cont'd)

- Row and key locks escalate to table locks, not page locks
- “Queries Decide the Isolation Level” (Erik Darling)

Optimized Locking

- Changes to the locking mechanisms
 - Locks on the Transaction ID
 - Releases locks on other objects quickly
 - Minimizes lock escalation
- Accelerated Database Recovery is Required

User DB

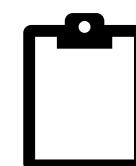
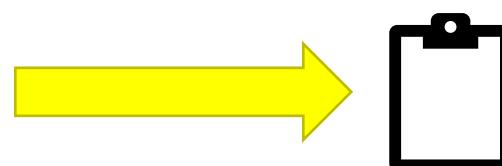


BEGIN TRAN 1A

```
UPDATE Customer  
SET State = 'MA'  
WHERE FirstName = 'Sebastian'
```



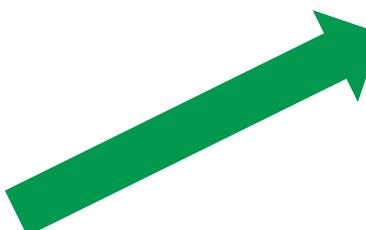
Sebastian	MA	Key	Padlock
-----------	----	-----	---------



Persistent Version Store

Sebastian	AZ	1A	1
-----------	----	----	---

```
SELECT FirstName,  
LastName,  
State  
FROM Customer  
WHERE FirstName = 'Sebastian'
```



Optimized Locking -
Reads

**Yeah, but that was just
a SELECT statement...**

Regular RCSI stuff.

Now what about writes?



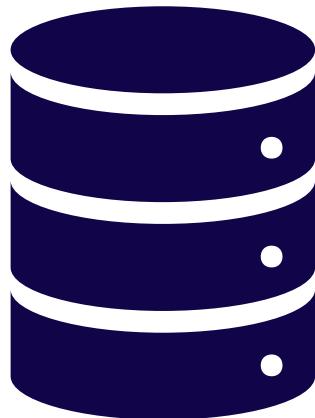
ON TOUR | NEW YORK

 **APASS**

Transaction ID Locking

- As the lock is held on the Transaction ID, the next transaction can see when rows aren't affected and update them without waiting.

User DB

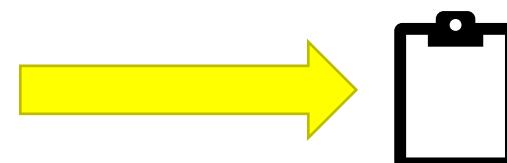


BEGIN TRAN 1A

```
UPDATE Customer  
SET State = 'MA'  
WHERE FirstName = 'Sebastian'
```



Sebastian	MA	Locked (Orange)	Locked (Orange)
Deborah	TN	Locked (Green)	Locked (Green)

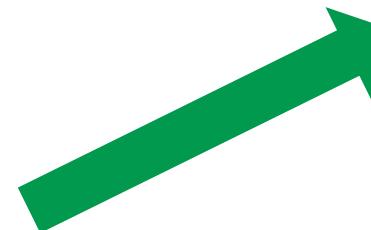


Persistent Version Store

Sebastian	AZ	1A	1
Deborah	TN	2A	1

BEGIN TRAN 2A

```
UPDATE Customer  
SET State = 'MA'  
WHERE FirstName = 'Deborah'
```

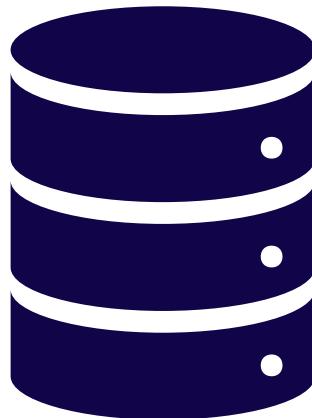


Optimized Locking -
TID Locking

Lock After Qualification (LAQ)

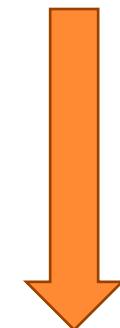
- RCSI is a requirement for this to work
- Looks to see if the previously committed version of the rows with TID locks would also be affected by the current transaction
 - If yes, transaction ***will wait*** for previous transaction to finish and include those columns as part of the change
 - If no, the transaction will update the rows that do match

User DB



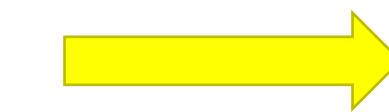
BEGIN TRAN 1A

```
UPDATE Customer  
SET State = 'GA'  
WHERE FirstName = 'Sebastian'
```



COMMIT TRAN 1A

Sebastian	MA	Key	Locked
-----------	----	-----	--------



BEGIN TRAN 2A

```
UPDATE Customer  
SET State = 'MA'  
WHERE FirstName = 'Sebastian'
```



Persistent Version Store				
FirstName	State	Version	TransID	Order
Sebastian	AZ	1A	1	
Sebastian	GA	2A	1	

Optimized Locking -
LAQ

Demos!



ON TOUR | NEW YORK

APASS

Be aware of

- Persistent Version Store
 - Make sure there is enough space available
 - When full, only reads can happen – no writes
 - Asynchronous Cleaner
 - Data remains until it's no longer needed
 - Keep Transactions Fast

Be aware of

- Lock After Qualification
 - Blocking and waits can still occur
 - Be aware of multiple transactions that may change data for the same set of records

Recap



- Understand how RCSI & ADR are used for optimized locking
- Minimizes blocking but does not solve ALL situations
- Reduces the number of locks being held and the length they are held for
- Writes don't block Reads OR Writes

Resources – Optimized Locking

- [SQL Server Transaction Locking and Row Versioning Guide](#)
- [Optimized Locking](#)
- [Bob Ward's Github - Optimized Locking Demos](#)
- [Isolation Level Locking \(Erik Darling\)](#)

Resources - Accelerated Database Recovery

- [Soaring to New Heights with Accelerated Database Recovery \(John Morehouse\)](#)
- [Monitor and Troubleshoot Accelerated Database Recovery](#)
- [Why to Use Accelerated Database Recovery in SQL Server \(Luis Lema\)](#)
- [Constant Time Recovery in Azure SQL Database \(whitepaper\)](#)

Your feedback is important to us



Evaluate this session at:

www.PASSDataCommunitySummit.com/evaluation

Thank you

Any Questions?

Deborah Melkin



[@dgmelkin.bsky.social](https://bsky.social/@dgmelkin)



github.com/DebtheDBA



[DebtheDBA.wordpress.com](https://debthedba.wordpress.com)

