

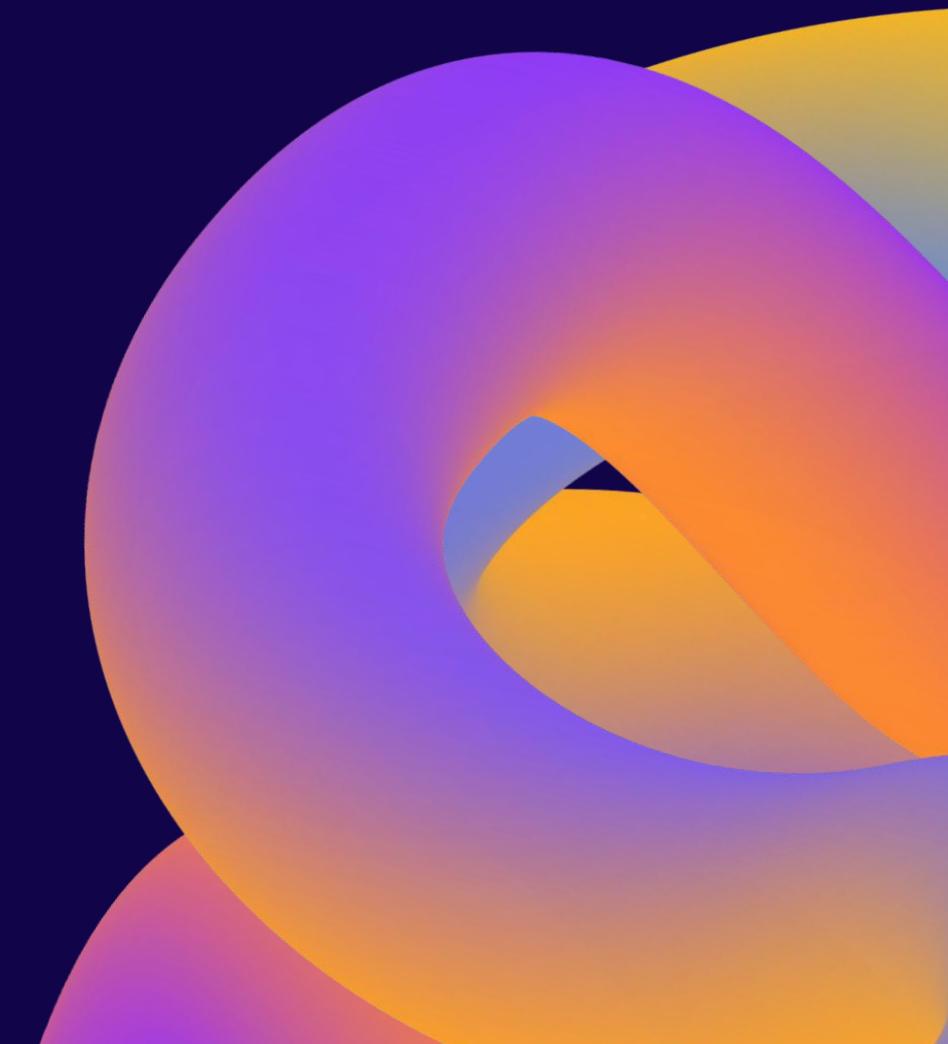
Optimized Locking: Improving SQL Server Transaction Concurrency

Deborah Melkin

she\her

Data Engineer

Advisor360



Deborah Melkin

She/Her

Data Engineer
Advisor360



- . 25 years as a DBA
- . Data Platform Women in Tech (WIT) Virtual UG, Co-leader
- . WITspiration, Co-founder
- . Redgate Community Ambassador
- . Microsoft MVP, Data Platform
- . In my spare time, I can usually be found doing something musical or something geeky with my husband, Andy, and our dog, Sebastian.



@dgmelkin.bsky.social



github.com/DebtheDBA



DebtheDBA.wordpress.com

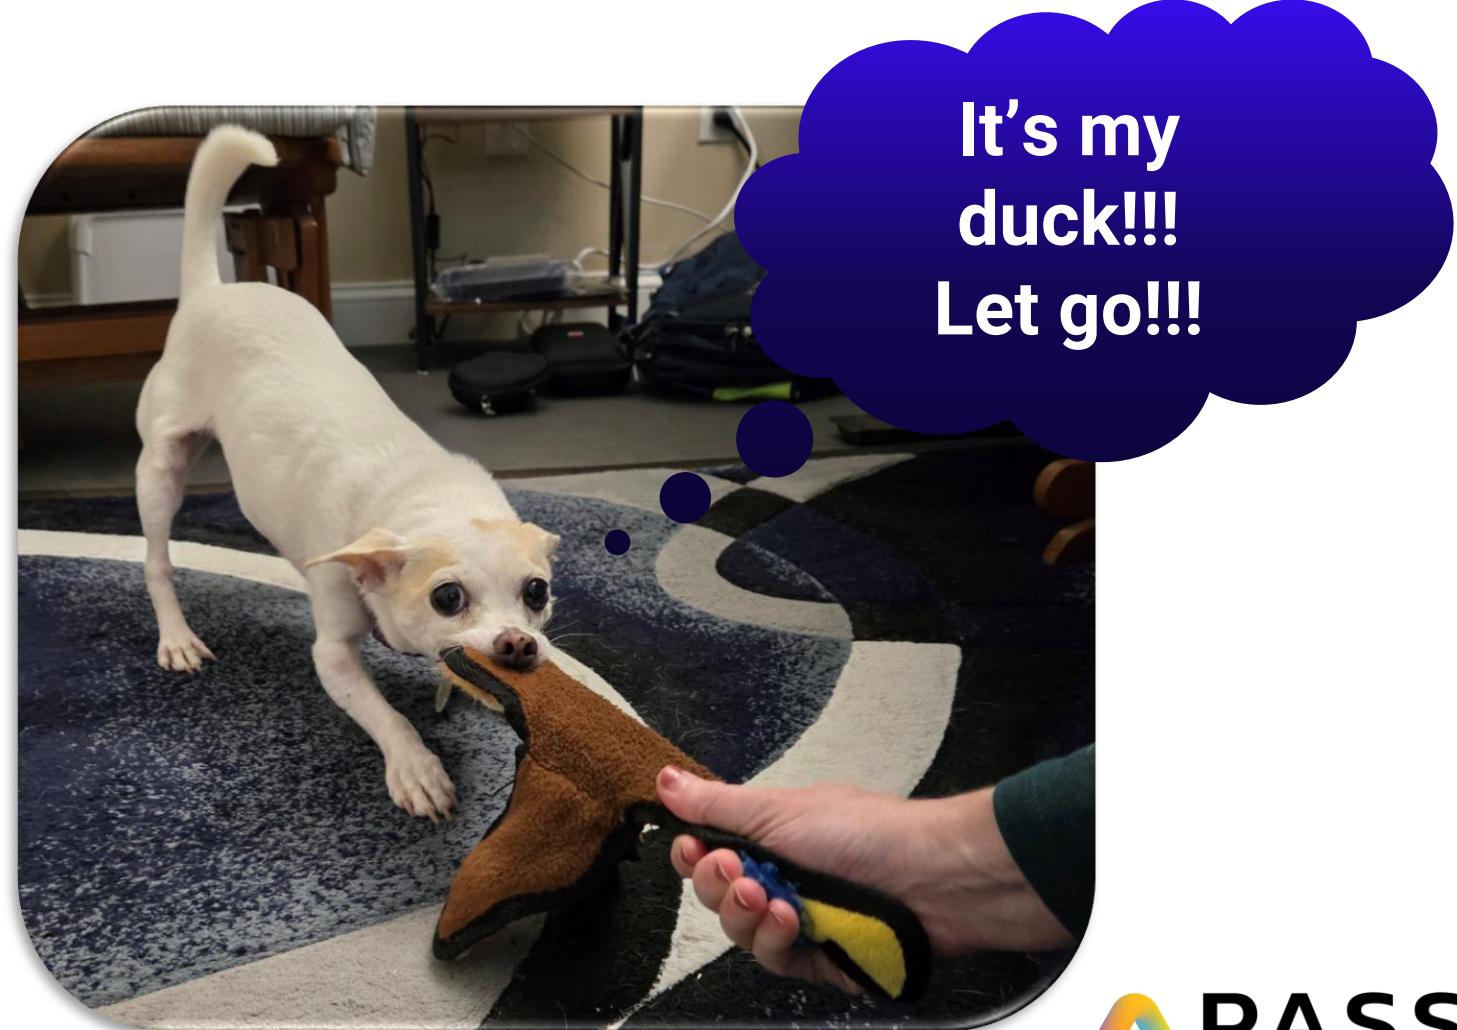
Hidden Pathways to Achieving Peak SQL Server Performance

- Revealing Hidden Performance Issues with Extended Events (Jeff Iannucci)
- Maximizing SQL Server Performance with Read Committed Snapshot Isolation (Haripriya Naidu)
- Optimized Locking: Improving SQL Server Transaction Concurrency (Deborah Melkin)
- Mitigating Your Data Bloat with Partitioning & Data Virtualization (Andy Yun)
- Deep Dive into Memory Optimized TempDB (Haripriya Naidu)



What's the problem we're trying to solve?

- Locking & Blocking!

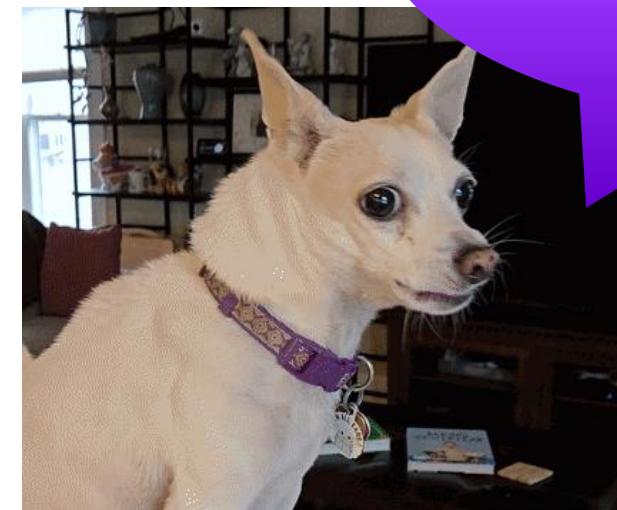


Introducing Optimized Locking

- GA in Azure SQL DB in Feb 2023
 - Enabled by default
- SQL Server 2025
 - Manually enable by database
- Writes not blocking Writes

Agenda – How it works

- Components used by Optimized Locking
 - Read Committed Snapshot Isolation Level (RCSI)
 - Accelerated Database Recovery (ADR)
- Lock Escalation
- Optimized Locking
 - Transaction ID Locking
 - Lock After Qualification



Warning:
*** Internals Ahead ***
Level 300

Pessimistic vs Optimistic Locking

- **Pessimistic Locking** – Preventing users from modifying data in a way that affects other users
 - Read Uncommitted
 - Read Committed
 - Serializable
 - Repeatable Read

Pessimistic vs Optimistic Locking

- **Optimistic Locking** – No locks when reading data but data is checked when modifying
 - Read Committed Snapshot Isolation
 - Snapshot Isolation
 - Optimized Locking

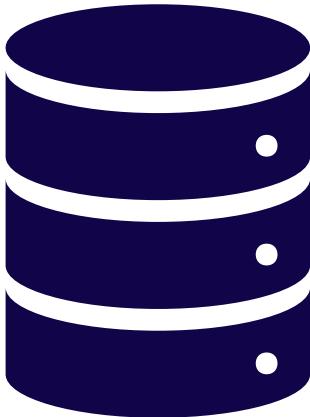
Read Committed Snapshot Isolation Level (RCSI)

- Introduced with SQL Server 2005
- Writes and Reads don't block each other
- Two database options for setting this:
 - SET READ COMMITTED SNAPSHOT ISOLATION – sets default for the database
 - ALLOW SNAPSHOT ISOLATION LEVEL – allows individual sessions\queries to use SNAPSHOT isolation levels even if read committed snapshot isolation is not set.

Read Committed Snapshot Isolation Level (RCSI)

- Every row in the database has a Transaction ID (TID)
- Keeps copies of the previously committed versions of the records in the Version Store in TempDB
- Only reading the committed version of the data

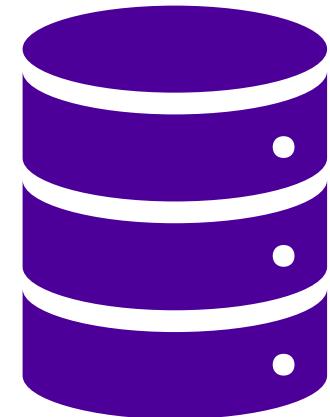
User DB



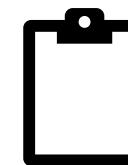
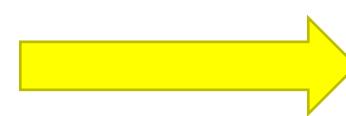
BEGIN TRAN 1A

```
UPDATE Customer  
SET State = 'MA'  
WHERE FirstName = 'Sebastian'
```

tempdb



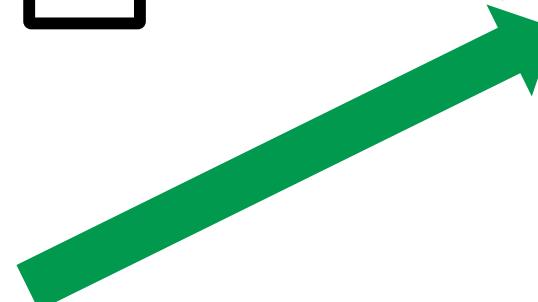
Sebastian	MA	Key	TID 1A
-----------	----	-----	-----------



Version Store			
Sebastian	AZ	TID 1A	1

TID
Transaction Seq No.

```
SELECT FirstName,  
LastName,  
State  
FROM Customer  
WHERE FirstName = 'Sebastian'
```



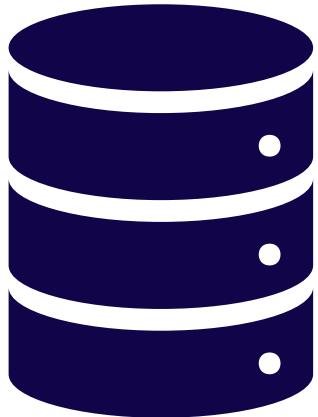
Read Committed
Snapshot Isolation

Accelerated Database Recovery (ADR)

- Introduced in SQL Server 2019
- Default for Azure SQL DB
 - Cannot be turned off
- Changes the way the transactions logs are read to be able to recover from long running transactions or just restore faster
- Uses a Persistent Version Store in Database rather than the Version Store in TempDB

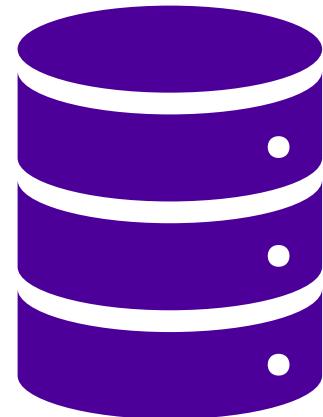
BEGIN TRAN 1A

User DB

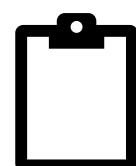


```
UPDATE Customer  
SET State = 'MA'  
WHERE FirstName = 'Sebastian'
```

tempdb



Sebastian	MA	Key	1A
-----------	----	-----	----



Persistent Version Store

Sebastian	AZ	1A	1
-----------	----	----	---

```
SELECT FirstName,  
LastName,  
State  
FROM Customer  
WHERE FirstName = 'Sebastian'
```

Version Store

Accelerated Database
Recovery

Lock Escalation

- Threshold at which SQL Server will convert lower granularity locks to a higher level in order to manage a small number of locks
 - 5,000 on a table or index for a single statement
 - Each additional 1,250 locks taken in a transaction

Lock Granularity (from Microsoft)

Resource	Description
RID	A row identifier used to lock a single row within a heap.
KEY	A row lock to lock a single row in a B-tree index.
PAGE	An 8 kilobyte (KB) page in a database, such as data or index pages.
EXTENT	A contiguous group of eight pages, such as data or index pages.
HoBT	A heap or B-tree. A lock protecting a B-tree (index) or the heap data pages in a table that doesn't have a clustered index.
TABLE	The entire table, including all data and indexes.
FILE	A database file.
APPLICATION	An application-specified resource.
METADATA	Metadata locks.
ALLOCATION_UNIT	An allocation unit.
DATABASE	The entire database.
XACT	Transaction ID (TID) lock used in Optimized Locking

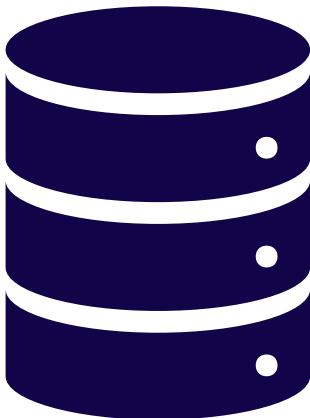
Lock Escalation (cont'd)

- Row and key locks escalate to table locks, not page locks
- “Queries Decide the Isolation Level” (Erik Darling)

Optimized Locking

- Changes to the locking mechanisms
 - Locks on the Transaction ID
 - Releases locks on other objects quickly
 - Minimizes lock escalation
- Accelerated Database Recovery is Required

User DB

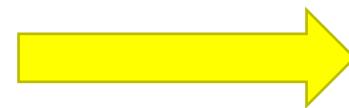
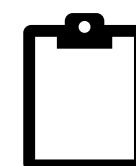
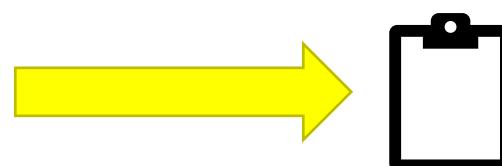


BEGIN TRAN 1A

```
UPDATE Customer  
SET State = 'MA'  
WHERE FirstName = 'Sebastian'
```



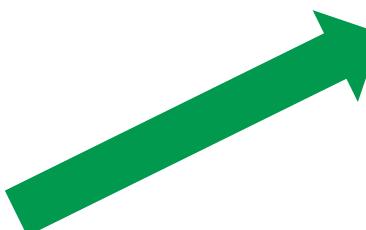
Sebastian	MA	Key	Padlock
-----------	----	-----	---------



Persistent Version Store

Sebastian	AZ	1A	1
-----------	----	----	---

```
SELECT FirstName,  
LastName,  
State  
FROM Customer  
WHERE FirstName = 'Sebastian'
```



Optimized Locking -
Reads

**Yeah, but that was just
a SELECT statement...**

Regular RCSI stuff.

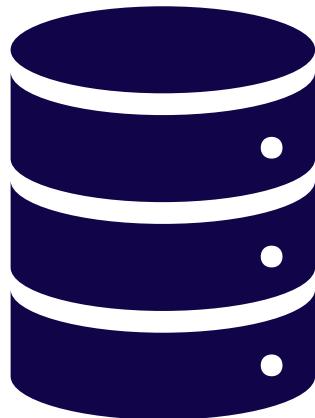
Now what about writes?



Transaction ID Locking

- As the lock is held on the Transaction ID, the next transaction can see when rows aren't affected and update them without waiting.

User DB

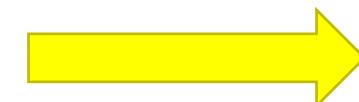
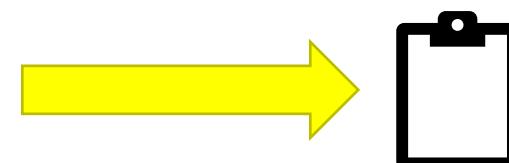


BEGIN TRAN 1A

```
UPDATE Customer  
SET State = 'MA'  
WHERE FirstName = 'Sebastian'
```



Sebastian	MA	Locked (Orange)	Locked (Orange)
Deborah	TN	Locked (Green)	Locked (Green)

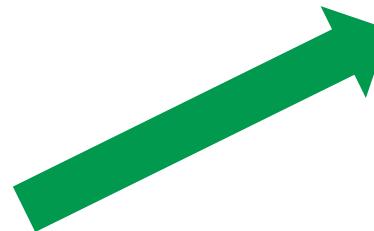


Persistent Version Store

Sebastian	AZ	1A	1
Deborah	TN	2A	1

BEGIN TRAN 2A

```
UPDATE Customer  
SET State = 'MA'  
WHERE FirstName = 'Deborah'
```

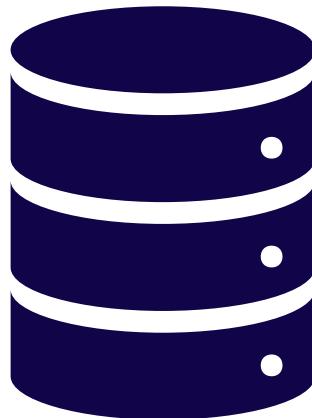


Optimized Locking -
TID Locking

Lock After Qualification (LAQ)

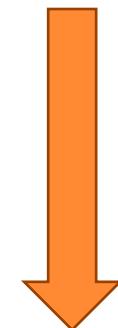
- RCSI is a requirement for this to work
- Looks to see if the previously committed version of the rows with TID locks would also be affected by the current transaction
 - If yes, transaction ***will wait*** for previous transaction to finish and include those columns as part of the change
 - If no, the transaction will update the rows that do match

User DB



BEGIN TRAN 1A

```
UPDATE Customer  
SET State = 'GA'  
WHERE FirstName = 'Sebastian'
```



COMMIT TRAN 1A

Sebastian	MA	Key	Locked
-----------	----	-----	--------



BEGIN TRAN 2A

```
UPDATE Customer  
SET State = 'MA'  
WHERE FirstName = 'Sebastian'
```



Persistent Version Store				
FirstName	State	Version	TransID	Order
Sebastian	AZ	1A	1	
Sebastian	GA	2A	1	

Optimized Locking -
LAQ

Additional Improvements with LAQ

- Announced with GA release of SQL Server 2025:
 - Skip index locks (SIL)
 - If there are no other queries that are requiring row locking queries (RLQ) for the rows needed, such as REPEATABLE READ or SERIALIZABLE isolations or queries with those corresponding hints), the engine can skip row and page level locks when modifying a row and just take an exclusive page lock

Additional Improvements with LAQ

- Announced with GA release of SQL Server 2025:
 - Query Plan LAQ Feedback Persistence
 - The qualification process may need to restart if there are other transactions modify the row before the current transaction, creating overhead. This was previously monitored on database level and now can be monitored on query level.
 - Previously would disable LAQ on database level if overhead exceeded a threshold. Can now be done on query level.

Demos!



Be aware of

- Persistent Version Store
 - Make sure there is enough space available
 - When full, only reads can happen – no writes
 - Asynchronous Cleaner
 - Data remains until it's no longer needed
 - **And** there are no more open transactions in front of it anywhere in the instance.
 - Keep Transactions Fast

Be aware of

- Lock After Qualification
 - Blocking and waits can still occur
 - Be aware of multiple transactions that may change data for the same set of records
 - Query Plan LAQ Feedback Persistance

Recap

- Understand how RCSI & ADR are used for optimized locking
- Minimizes blocking but does not solve ALL situations
- Reduces the number of locks being held and the length they are held for
- Writes don't block Reads OR Writes



Resources – Optimized Locking

- [SQL Server Transaction Locking and Row Versioning Guide](#)
- [Optimized Locking](#)
 - [Introducing Optimized Locking V2](#)
- [Bob Ward's Github - Optimized Locking Demos](#)
- [Isolation Level Locking \(Erik Darling\)](#)

Resources - Accelerated Database Recovery

- [Soaring to New Heights with Accelerated Database Recovery \(John Morehouse\)](#)
- [Monitor and Troubleshoot Accelerated Database Recovery](#)
- [Why to Use Accelerated Database Recovery in SQL Server \(Luis Lema\)](#)
- [Constant Time Recovery in Azure SQL Database \(whitepaper\)](#)

Thank you

Any Questions?

Deborah Melkin



[@dgmelkin.bsky.social](https://dgmelkin.bsky.social)



github.com/DebtheDBA



DebtheDBA.wordpress.com



Your feedback is important to us



Evaluate this session at:

passdatacommunitysummit.com/evaluations